


Learning to Boost Training by Periodic Nowcasting Near Future Weights

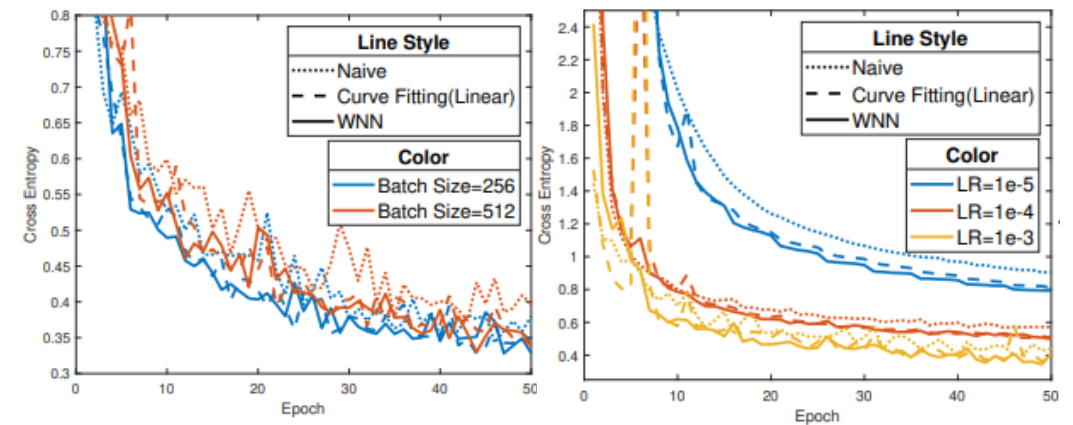
Jinhyeok Jang, Woo-han Yun, Won Hwa Kim,
Youngwoo Yoon, Jaehong Kim, Jaeyeon Lee,
ByungOk Han

Weight Prediction

- Can we bypass the training process and directly arrive at future weights?



Description of Optimization as Contour



General Tendency of Loss graph

Our Strategy

- 1) Learning-based regression model
- 2) Element-wise independent forecasting.
- 3) Separate forecasting for each mathematical operation,
- 4) Periodic short-term nowcasting per every 5 epochs,
- 5) Predicting residual between the future and the current weights
- 6) Applying a forecast network trained on Image Classification Task to various tasks




Figure 4. Conceptual view of short-term prediction with the proposed WNN.

Weight Nowcaster Network (WNN)

- For weight forecast, two things were considered:
 - 1) **Accurate Regression**
 - 2) **Fast Process**: $\sim 5,000$ parameters size of WNN




Figure 5. **Weight Nowcaster Network Architecture.** The WNN is composed of simple two-stream networks that use fully-connected layers and an activation network. Feature vectors from those two networks are unified to a feature vector and it is passed through a fully-connected layer. The predicted future weight parameters are obtained by adding outputs and input weight parameters.

Data Collection

- Architecture: LeNet, VGG, ResNet, MobileNetV2, ShuffleNetV2, DenseNet
- Dataset: MNIST, CIFAR10
- Optimizer: Adam
- Training Data: Weights of 30,000 epochs ($\sim 1.8e+10$ parameters)
- Validation: 2,200 epochs of ShuffleNetV2 and ResNet32

Experiments I

- Application to Vanilla CNN + CIFAR10
- Comparisons with
 - 1) Various Curve Fitting Models
 - 2) Various Acceleration Methods

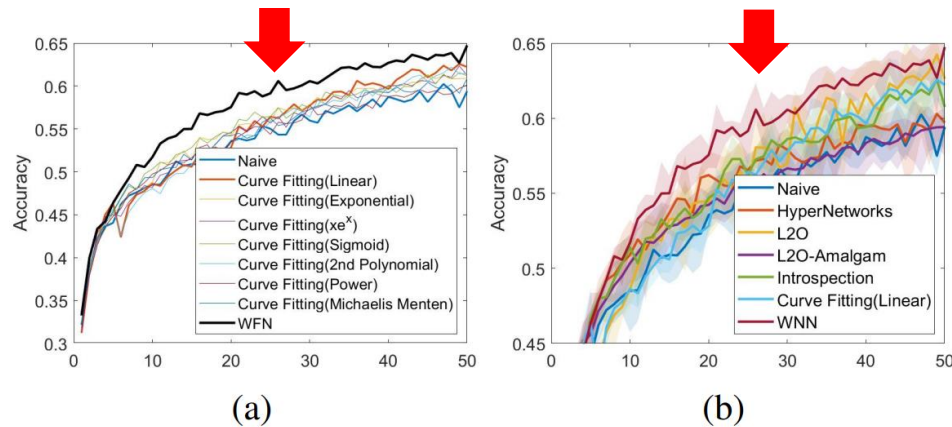


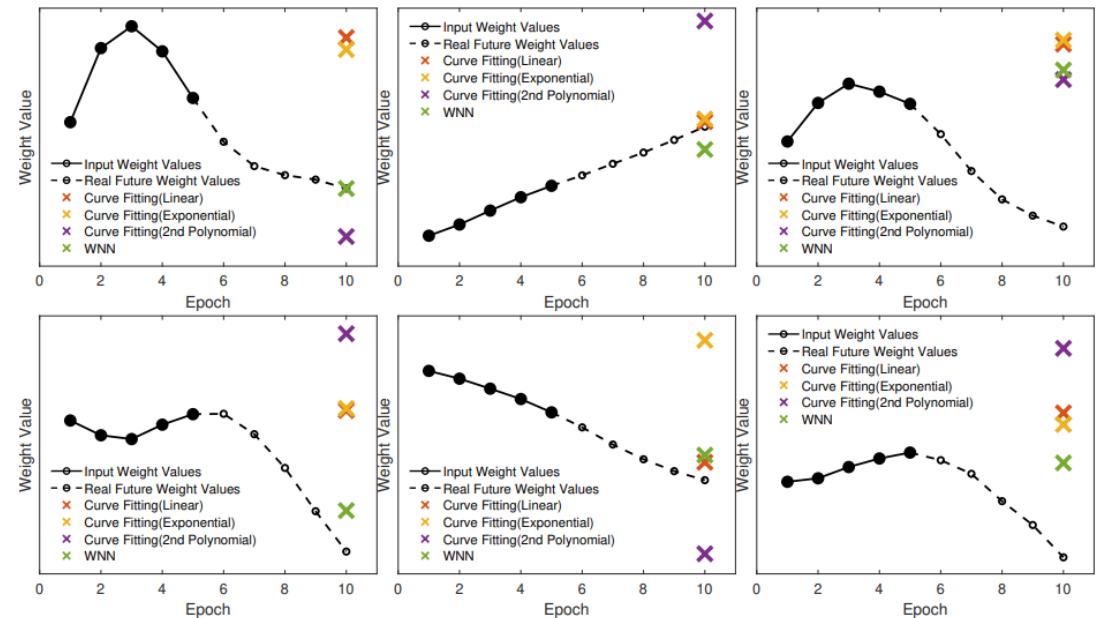

Figure 5. Comparisons of validation accuracy (a) of the proposed and curve fitting models, (b) of the proposed with previous methods (HyperNetworks, L2O, L2O-Amalgam, and Introspection). The shading represents the variation in validation accuracy of five trials.

Table 4. Time cost comparisons to train a Vanilla CNN by using various recent methods on the CIFAR10 dataset. NVIDIA TITAN Xp GPU was used to estimate time cost. “Converge” is the time to reach a validation accuracy of 59%

Method	Update (sec/batch)	meta-learning (hour)	forecasting (sec)	Converge (sec)	Speed Up
Naive Training	0.0245	-	-	52.82	$\times 1.00$
HyperNetworks	0.0267	-	-	51.02	$\times 1.04$
L2O	0.0290	20 (per task)	-	44.05	$\times 1.20$
L2O-Amalgam	0.0291	6.35 (only once)	-	67.02	$\times 0.79$
Introspection	0.0245	0.03 (only once)	0.015	43.23	$\times 1.22$
CF (Linear)	0.0245	-	0.015	39.71	$\times 1.33$
CF (Exponential)	0.0245	-	3.32	61.94	$\times 0.85$
WNN	0.0245	0.08 (only once)	0.015	25.27	$\times 2.09$

Experiments II

- Application to Vanilla CNN and CIFAR10
- Comparisons with Linear Curve Fitting Model



Experiments III

• Further Tasks

- (i) **ImageNet Classification.** ImageNet (Deng et al., 2009) is a widely-used large-scale dataset with 1.3 million images from 1,000 classes. We trained the MobileNetV2 using the Adam with exponential decay.
- (ii) **Image Segmentation.** DeepLabV3+ (Chen et al., 2018) with the MobileNetV2 backbone was trained on the PASCAL VOC 2012 dataset (Everingham et al., 2015) with the SGD, cosine decay, and cross entropy.
- (iii) **Pose Estimation.** We adopted OpenPose (Cao et al., 2017) with an ImageNet-pre-trained VGG19 backbone on the MS COCO 2016 (Lin et al., 2014) that consists of over 100K persons' keypoints.
- (iv) **Language Modeling.** The goal is to predict a masked word from a sequence of words, *i.e.*, text. The universal BERT (Dehghani et al., 2019) on the WikiText-2 dataset which consists of over 2 million words was trained with the Adam, the masked LM loss, and the penalized confidence (Pereyra et al., 2017).
- (v) **Reinforcement Learning.** We applied WNN to the Pendulum problem, which is a famous reinforcement learning problem using the DDPG (Deep Deterministic Policy Gradient) (Lillicrap et al., 2015). Two networks, an actor and a critic, were trained using the Adam for 200 episodes.
- (vi) **Transfer Learning on Attention Model.** PVTv2-B0 (Wang et al., 2022) with ImageNet-pre-trained weights was trained on CIFAR100 for 50 epochs using the Adam, the warm-up scheduling, and decay.
- (vii) **Diffusion Model.** We validated on the Denoising Diffusion Implicit Model (DDIM) (Song et al., 2020) on the Oxford Flowers dataset (Nilsback & Zisserman, 2008). The model was trained to minimize ℓ_1 loss using the AdamW with the learning rate decay from $1e-3$ for 60 epochs. We evaluated it based on KID (Kernel Inception Distance) metric (Binkowski et al., 2018).

Table 5. Experimental comparisons on various tasks with naive training, Introspection, a linear curve fitting, and the proposed WNN. WNN consistently outperforms the other methods on a variety of tasks.

Task (Metric)	Method	Best	Converge (epoch/episode)	Reach	Speed Up	Task (Metric)	Method	Best	Converge (epoch/episode)	Reach	Speed Up
ImageNet Classification (Val Acc \uparrow)	Naive	69.40%	26	68.00%	$\times 1.00$	RL (Episode Reward \uparrow)	Naive	-139.10	116	-200.00	$\times 1.00$
	Introspection	70.13%	24	68.00%	$\times 1.08$		Introspection	-128.70	107	-200.00	$\times 1.08$
	CF (Linear)	68.19%	41	68.00%	$\times 0.63$		CF (Linear)	-138.79	111	-200.00	$\times 1.05$
	WNN	70.57%	21	68.00%	$\times \mathbf{1.24}$		WNN	-123.55	88	-200.00	$\times \mathbf{1.32}$
Image Segmentation (Val Jaccard \uparrow)	Naive	53.99%	55	48.00%	$\times 1.00$	Transfer Learning on Attention Model (Val Acc \uparrow)	Naive	83.22%	17	80.00%	$\times 1.00$
	Introspection	53.97%	46	48.00%	$\times \mathbf{1.20}$		Introspection	81.52%	38	80.00%	$\times 0.45$
	CF (Linear)	53.79%	50	48.00%	$\times 1.10$		CF (Linear)	82.51%	14	80.00%	$\times \mathbf{1.21}$
	WNN	53.81%	45	48.00%	$\times \mathbf{1.20}$		WNN	83.09%	14	80.00%	$\times \mathbf{1.21}$
Pose Estimation (Val Loss \downarrow)	Naive	623.54	43	673.00	$\times 1.00$	Diffusion Model (Val KID \downarrow)	Naive	0.1918	32	0.2000	$\times 1.00$
	Introspection	627.17	46	673.00	$\times 0.94$		Introspection	0.1786	36	0.2000	$\times 0.89$
	CF (Linear)	614.80	30	673.00	$\times \mathbf{1.43}$		CF (Linear)	0.1878	54	0.2000	$\times 0.59$
	WNN	615.79	30	673.00	$\times \mathbf{1.43}$		WNN	0.1732	26	0.2000	$\times \mathbf{1.23}$
Language Modeling (Val Perplexity \downarrow)	Naive	31.25	55	100.00	$\times 1.00$	Average	Naive	N/A	N/A	N/A	$\times 1.00$
	Introspection	39.93	56	100.00	$\times 0.98$		Introspection	N/A	N/A	N/A	$\times 0.94$
	CF (Linear)	42.67	58	100.00	$\times 0.95$		CF (Linear)	N/A	N/A	N/A	$\times 0.99$
	WNN	31.26	48	100.00	$\times \mathbf{1.15}$		WNN	N/A	N/A	N/A	$\times \mathbf{1.25}$

Code Availability

- <https://github.com/jjh6297/WNN>

Weight Nowcasting Network (WNN)

Code for ["Learning to Boost Training by Periodic Nowcasting Near Future Weights"]

dependency

| Library | Known Working | Known Not Working |
| tensorflow | 2.3.0, 2.9.0 | <= 2.0 |

Usage

WNN can be easily used as a callback function extending `tf.keras.callbacks.Callback`:

```
import tensorflow as tf
import tensorflow.keras
from WNN import *

.
.
.

model.fit(..., callbacks=[WeightForecasting()])
```

Pre-trained Weights

Pre-trained weights of WNN are included. 'NWNX_XXX_13.h5' in this repo are the pre-trained weights for each mathematical operation type (Conv, FC, Bias).

Thanks